

CS 31007

Autumn 2021

# COMPUTER ORGANIZATION AND ARCHITECTURE

---

Instructors

Rajat Subhra Chakraborty (*RSC*)

Bhargab B. Bhattacharya (*BBB*)

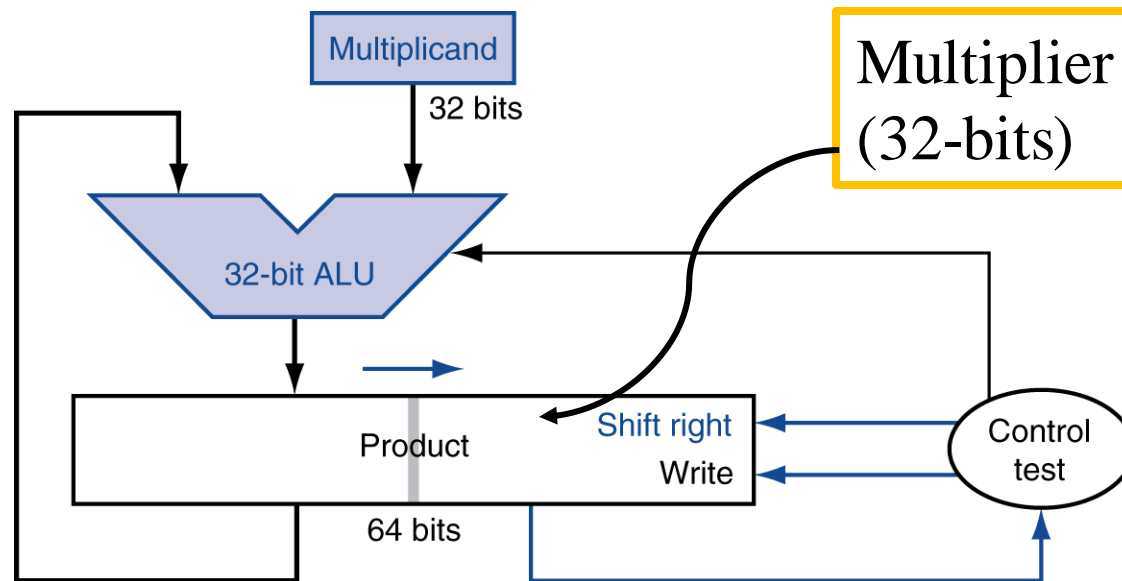
Lecture #22: Computer Arithmetic

20 September 2021

---

Indian Institute of Technology Kharagpur  
*Computer Science and Engineering*

# Recap: Integer Multiplier using Repeated Add/Shift



- Can we expedite multiplication?
- Handling negative numbers?
- - Booth's Algorithm

# Recall: Multiplication Example

$0010 \times 0110 = ?$   $0010$  (+ 2, multiplicand);  $0110$  (+ 6, multiplier)

Iteration	multiplier	Original algorithm	
		Step	Product
0	0010	Initial values	0000 0110
1	0010	1:0 $\Rightarrow$ no operation	0000 0110
	0010	2: Shift right Product <b>logical shift</b>	0000 0011
2	0010	1a:1 $\Rightarrow$ prod = Prod + Mcand	0010 0011
	0010	2: Shift right Product	0001 0001
3	0010	1a:1 $\Rightarrow$ prod = Prod + Mcand	0011 0001
	0010	2: Shift right Product	0001 1000
4	0010	1:0 $\Rightarrow$ no operation	0001 1000
	0010	2: Shift right Product <b>+12</b>	0000 1100

# Booth's Encoding

(valid for signed multiplication as well)

- Recall old trick

Example:  $123454 \times 9$

$\Rightarrow$  six partial products plus addition of six numbers

- $123454 \times 9 = 123454 \times (10 - 1) = 1234540 - 123454$
  - Transform **addition of six partial products** to
  - **one shift and one subtraction!**
- Booth's algorithm applies the same principle
    - in binary we have just '1' and '0'

# Booth's Encoding

- Multiply  $x$  by 0111

$x$  (multiplicand); 0111 (multiplier)  $\Rightarrow x \times 7$

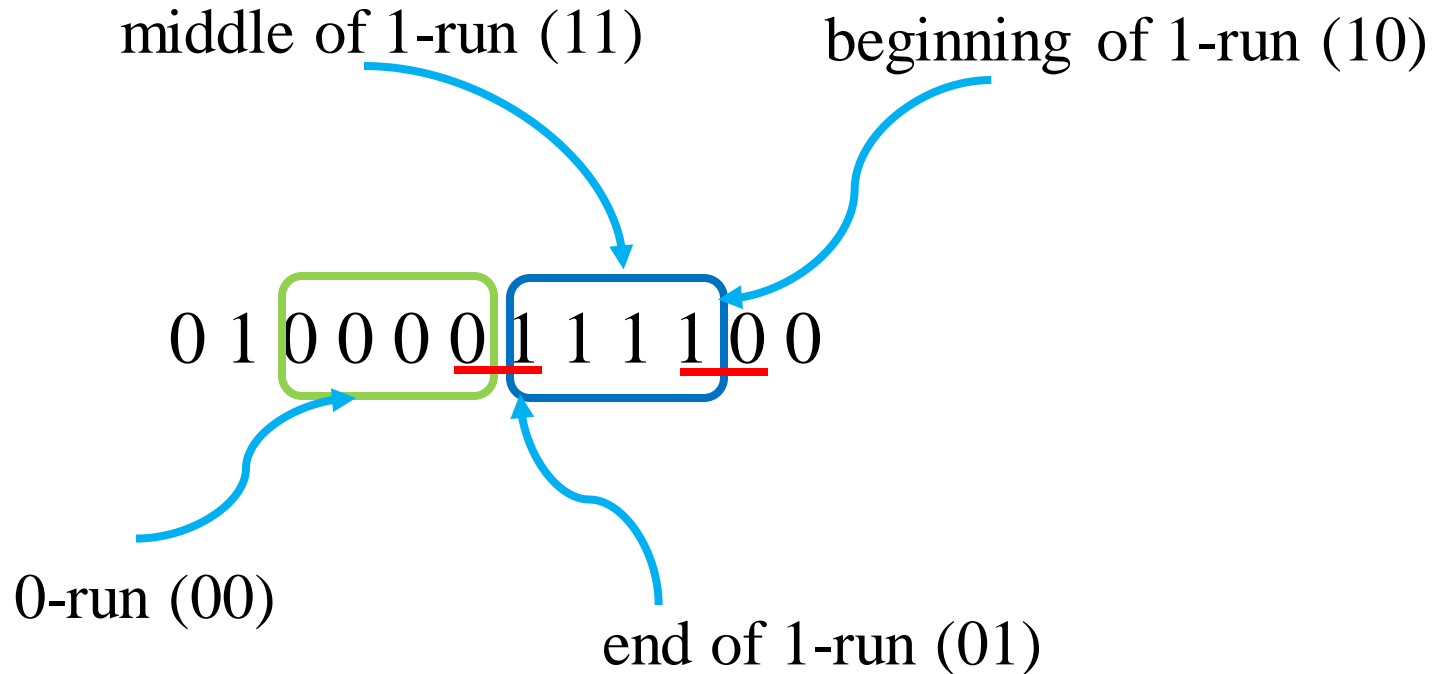
Search for a run of '1' bits in the multiplier

- e.g. '0111' has a run of 3 consecutive '1' bits
- Multiplying by '0111' (7 in decimal) is equivalent to multiplying by 8 and subtracting once, since

$$x \times 7 = x \times (8 - 1) = 8x - x \Rightarrow (\text{shift-left } x, \text{ three times}) - x$$

- Hence, iterate right to left and look for “runs of 1”:
- $x \times 111 = x \times (1000 - 1) = (x \times 2^3) - x$ 
  - Subtract multiplicand from product at first '1'
  - Shift multiplicand by 3-bits on left and add to the partial product after the last '1'
  - Do nothing for the consecutive 1-bits/in the middle, or for 0-runs (actually, we keep on shifting the product register)

# Booth's Encoding for Multiplier



# Booth's Algorithm

Current bit	Bit to right	Explanation	Example (multiplier)	Operation
1	0	Begins run of '1'	0000111 <b>1</b> 000	Subtract/shift
1	1	Middle of run of '1'	000011 <b>11</b> 000	Nothing/shift
0	1	End of a run of '1'	000 <b>01</b> 111000	Add/shift
0	0	Middle of a run of '0'	0 <b>00</b> 01111000	Nothing/shift

# Booth's algorithm: Example

$0010 \times 1101 = ?$   $0010$  (+ 2, multiplicand); **1101** (- 3, multiplier)

arithmetic shift

op performed in the left half

initialization

Iteration	multiplier	Booth's algorithm	
		Step	Product
0	0010	Initial values	<u>0000</u> 1101 0
1	0010	1c: 10 $\Rightarrow$ prod = Prod - Mcand	1110 1101 0
	0010	2: Shift right Product	1111 0110 1
2	0010	1b: 01 $\Rightarrow$ prod = Prod + Mcand	0001 0110 1
	0010	2: Shift right Product	0000 1011 0
3	0010	1c: 10 $\Rightarrow$ prod = Prod - Mcand	1110 1011 0
	0010	2: Shift right Product	1111 0101 1
4	0010	1d: 11 $\Rightarrow$ no operation	1111 0101 1
	0010	2: Shift right Product	1111 1010 1

start of 1-run

end of 1-run



# MIPS Multiplication

- Two 32-bit registers for product
  - HI: most-significant 32 bits
  - LO: least-significant 32-bits
- Instructions
  - `mult rs, rt` / `multu rs, rt`
    - 64-bit product in HI/LO
  - `mfhi rd` / `mflo rd`
    - Move from HI/LO to rd
    - Can test HI value to see if product overflows 32 bits
  - `mul rd, rs, rt`
    - Least-significant 32 bits of product → rd

# So far covered in computer arithmetic...

- ❖ Integer Number Systems and Overflow
  - ❖ Ripple-Carry Adder (RCA)
  - ❖ Carry-Lookahead Adder (CLA)
  - ❖ Hybrid Adder, CLT
  - ❖ Carry-Select Adder (CSA)
  - ❖ Brent-Kung's Parallel Prefix Adder (PPA)
  - ❖ Carry-Save Adders (for adding multiple operands)
  - ❖ Integer Multiplication
- 
- ❖ Integer Division: Reading Assignment
  - ❖ Floating-Point Arithmetic and Hardware

# Division

- Implemented by successive subtractions
- Result must verify the equality

$$\text{Dividend} = (\text{Multiplier} \times \text{Quotient}) + \text{Remainder}$$

Another powerful division method (Goldschmidt's algorithm):

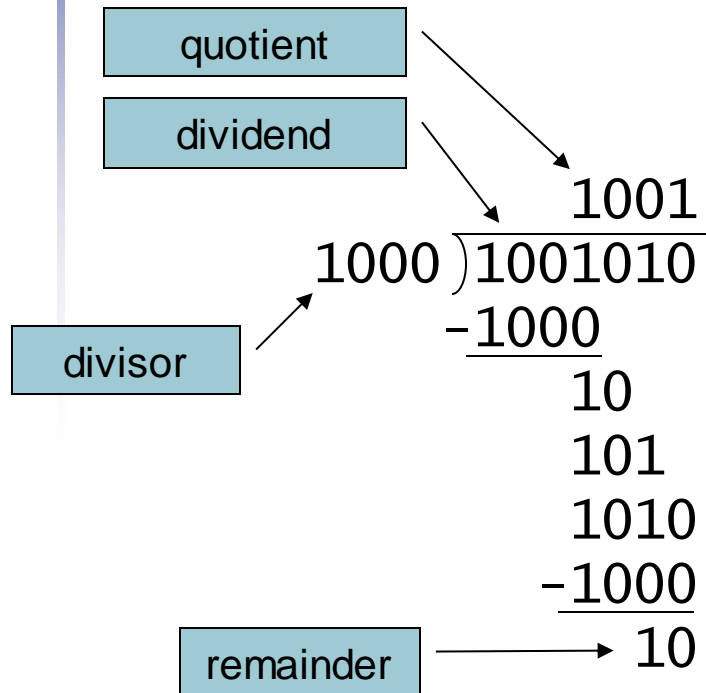
In computer science, division can be accomplished by multiplying a number with the reciprocal of the multiplier!

# Decimal division

$$\begin{array}{r} 303 \\ 7 \overline{) 2126} \\ \underline{-210} \phantom{0} \\ 26 \\ \underline{-21} \\ 5 \end{array}$$

- What are the rules?
  - Repeatedly try to subtract a multiple of divisor from dividend
  - Record multiple (or zero)
  - At each step, repeat with a lower power of ten
  - Stop when remainder is smaller than divisor

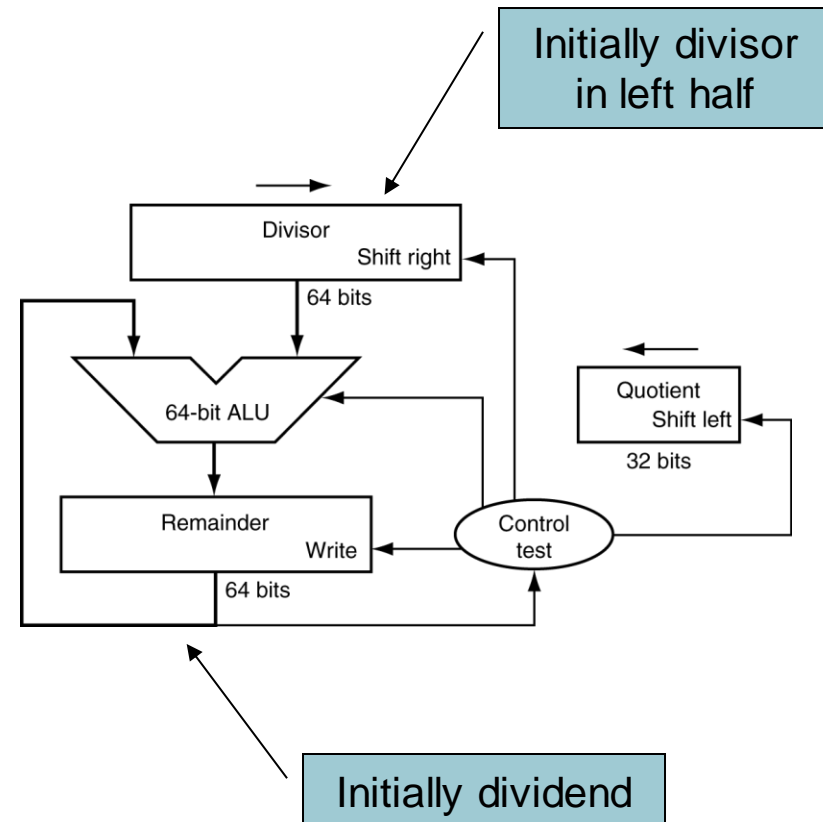
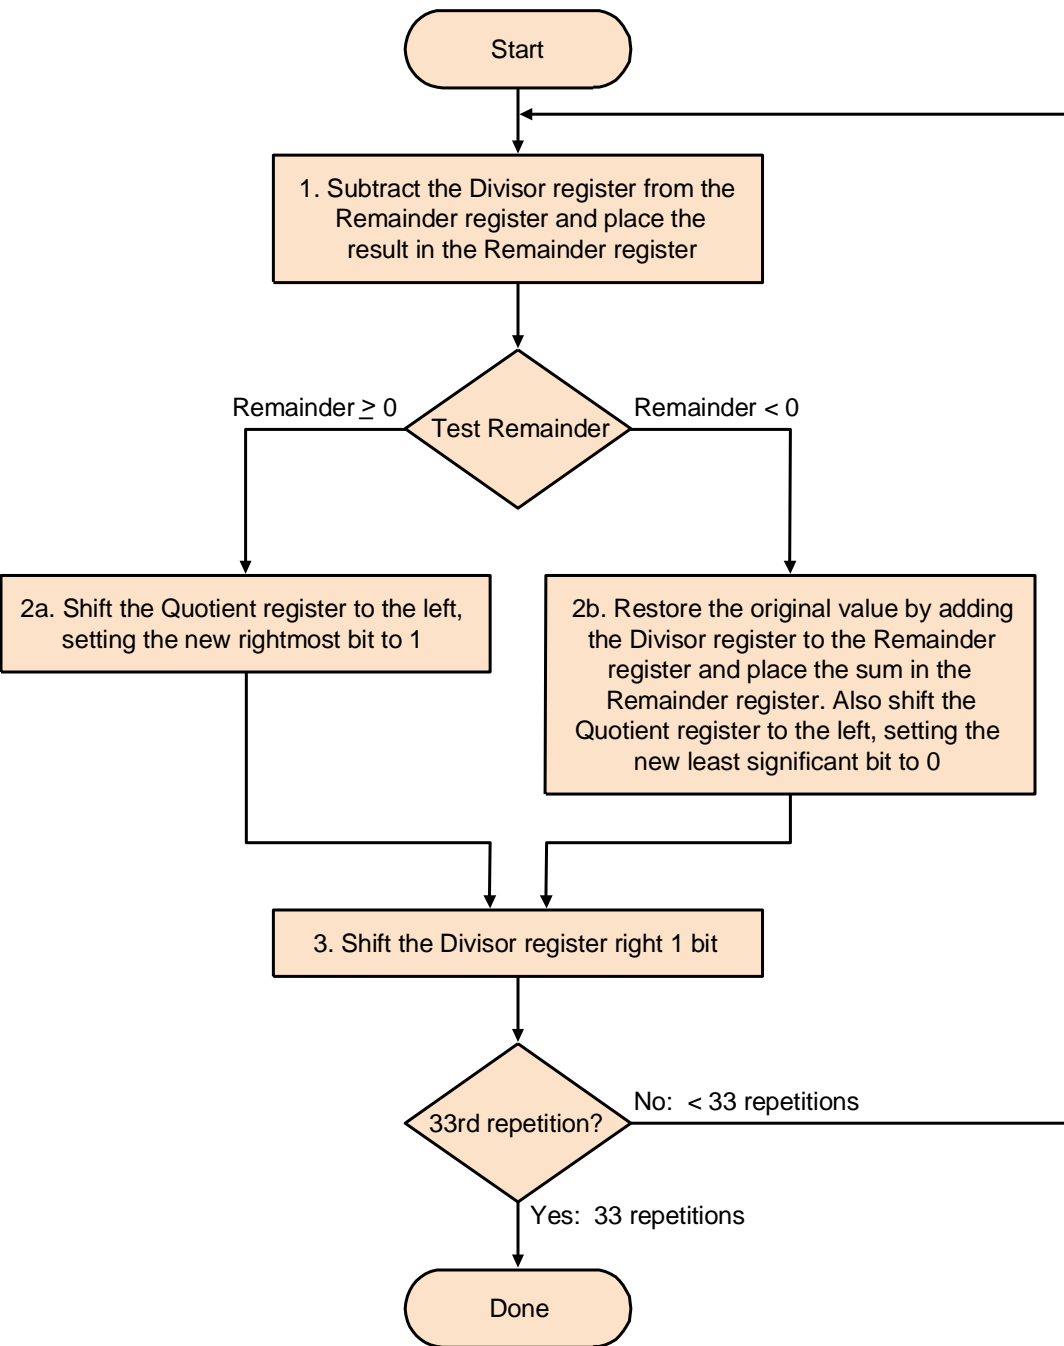
# Integer Division in Binary



$n$ -bit operands yield  $n$ -bit quotient and remainder

- Check for 0 divisor
- Long division approach
  - If divisor  $\leq$  dividend bits
    - 1 bit in quotient, subtract
  - Otherwise
    - 0 bit in quotient, bring down next dividend bit
- Restoring division
  - Do the subtract, and if remainder goes  $< 0$ , add divisor back
- Signed division
  - Divide using absolute values
  - Adjust sign of quotient and remainder as required

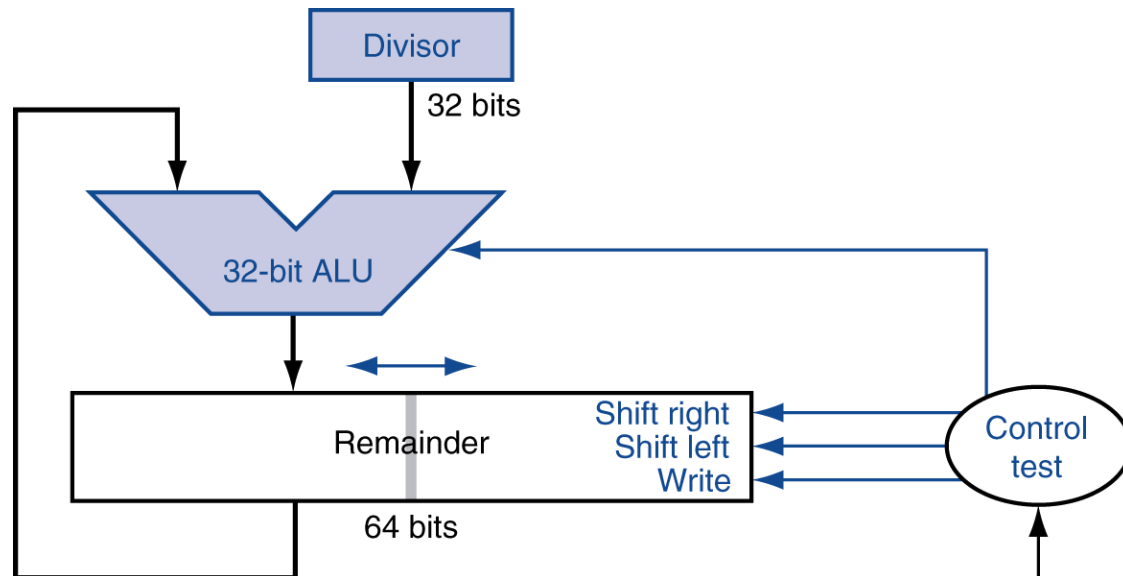
# Division Hardware



# Restoring Division

Iteration	Divisor	Divide algorithm	
		Step	Remainder
0	0010	Initial values	0000 0111
	0010	Shift Rem left 1	0000 1110
1	0010	2: Rem = Rem - Div	1110 1110
	0010	3b: Rem < 0 $\Rightarrow$ + Div, sll R, R0 = 0	0001 1100
2	0010	2: Rem = Rem - Div	1111 1100
	0010	3b: Rem < 0 $\Rightarrow$ + Div, sll R, R0 = 0	0011 1000
3	0010	2: Rem = Rem - Div	0001 1000
	0010	3a: Rem $\geq$ 0 $\Rightarrow$ sll R, R0 = 1	0011 0001
4	0010	2: Rem = Rem - Div	0001 0001
	0010	3a: Rem $\geq$ 0 $\Rightarrow$ sll R, R0 = 1	0010 0011
Done	0010	shift left half of Rem right 1	0001 0011

# Optimized Divider



- One cycle per partial-remainder subtraction
- Looks a lot like a multiplier!
  - Same hardware can be used for both



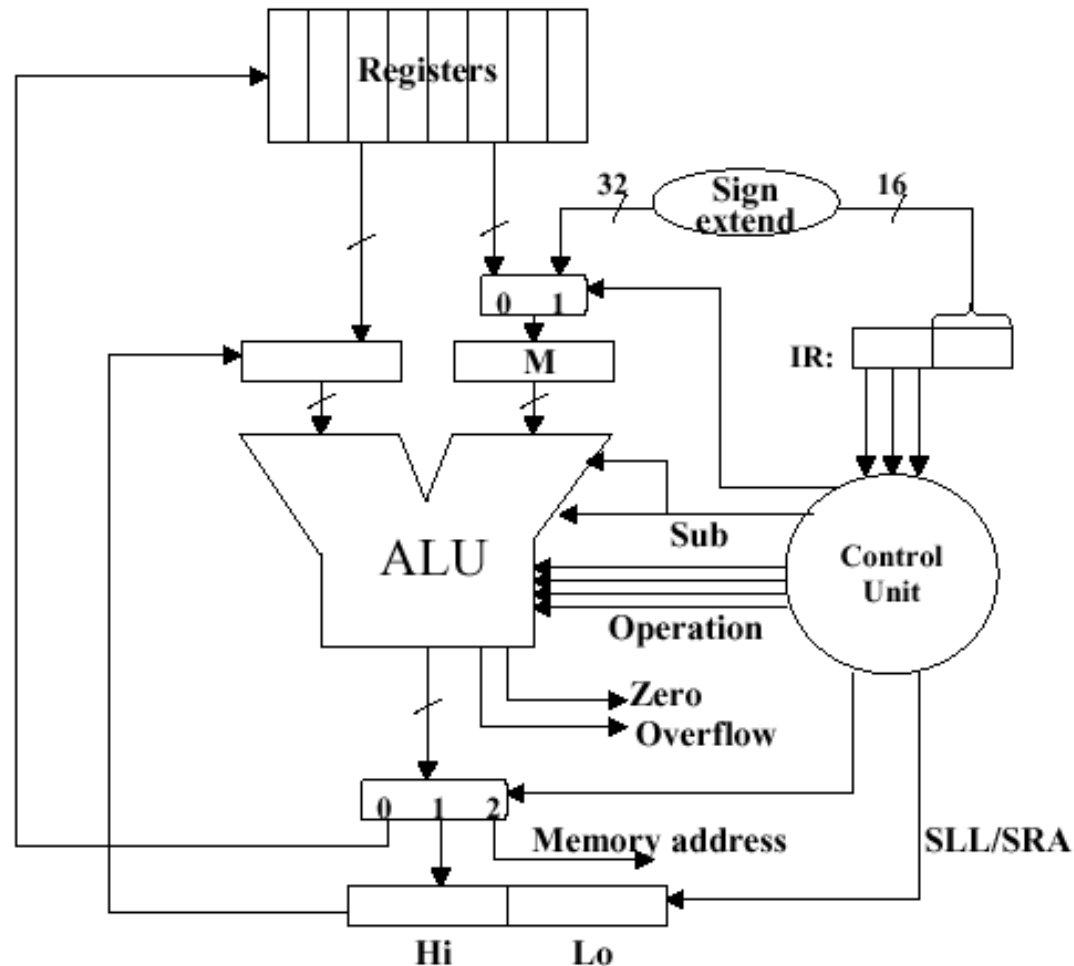
# Faster Division

- Can't use parallel hardware as in multiplier
  - Subtraction is conditional on sign of remainder
- Faster dividers (e.g. SRT division) generate multiple quotient bits per step
  - Still require multiple steps

# MIPS Division

- Use HI/LO registers for result
  - HI: 32-bit remainder
  - LO: 32-bit quotient
- Instructions
  - `div rs, rt` / `divu rs, rt`
  - No overflow or divide-by-0 checking
    - Software must perform checks if required
  - Use `mfhi`, `mflo` to access result

# MIPS Architecture for Integer Arithmetic: Multiplication and Division



CS 31007

Autumn 2021

# COMPUTER ORGANIZATION AND ARCHITECTURE

---

Instructors

Rajat Subhra Chakraborty (*RSC*)

Bhargab B. Bhattacharya (*BBB*)


Lecture #23, #24: Floating-Point Arithmetic

21 September 2021



---

Indian Institute of Technology Kharagpur  
*Computer Science and Engineering*

# Floating Point: Format, Arithmetic, and Hardware Implementation



A fundamental theoretical question: Can we prove Kepler's Law of planetary motion by virtue of experiments?



## The IEEE Floating-Point Standard

So far as the theories of mathematics are about reality, they are not certain; so far as they are certain, they are not about reality. - *Albert Einstein*

# A computation error observed by a UG student led to the ACM Turing Award later ....

- 1953: Willian Kahan, a UG student of Math at the *University of Toronto* was simulating numerically the dynamics of the wing controller of an aircraft during take-off and landing
- Observed certain unexpected results due to errors in computation
- => concept of floating-point (FP) arithmetic
- => principal architect behind **IEEE 754 FP standard (1985)**
- => Kahan honored with ACM Turing Award (1989)



William Kahan  
(1933 - )

# Floating-Point Representation

- Used to represent *real numbers*:  $-34.986 \times 10^{-22}$ ,  $\pi$ ,  $e$ ,  $\sqrt{2}$
  - Defined by IEEE 754 Standard
    - Kahan (1985)
  - Developed in response to divergence of representations
    - Portability issues for scientific code
  - Now almost universally adopted
  - Two representations
    - Single precision (32-bit)
    - Double precision (64-bit)
- Infamous Intel Pentium Bug (1994) => FDIV =>  
loss of \$300 Million




# IEEE Floating-Point Format

single: 8 bits

double: 11 bits

single: 23 bits

double: 52 bits



Fraction  
≡ Mantissa  
≡ Significand


$$N = (-1)^S \times (1.M) \times 2^{(\text{Exponent} - \text{Bias})}$$

- S: sign bit (0  $\Rightarrow$  non-negative, 1  $\Rightarrow$  negative)
- Normalize significand:  $1.0 \leq |\text{significand}| < 2.0$ 
  - Always has a leading pre-binary-point 1 bit, **so no need to represent it explicitly (hidden bit)**
  - Normalized Significand is Mantissa with the **“1.” restored**
- Exponent: excess representation: actual exponent + Bias
  - Ensures exponent is unsigned
  - Single: Bias = 127; Double: Bias = 1023
  - Trade-off between range (E) and precision (M)

# IEEE 754 Single-Precision Floating-Point

Representation of floating point numbers in IEEE 754 standard:

single precision



*FP-exponent E:*

excess 127 binary coding

Actual exponent is  $e = E - 127$

Floating-point exponent  $E = e + 127$

*mantissa:*

sign + magnitude, normalized  
binary significand w/ hidden  
integer bit: 1.M

FP-exponent ( $E$ )	Actual exponent $e$
<b>00000000 (0)</b>	<b>Special use</b>
00000001 (1)	- 126
00000010 (2)	- 125
.....	.....
<b>01111111 (127)</b>	<b>0</b>
.....	.....
11111110 (254)	+ 127
<b>11111111 (255)</b>	<b>Special use</b>

# IEEE 754 floating-point standard

- Leading “1” bit of significand is implicit
- Exponent is “biased” to make comparison easier
  - all 0s is smallest exponent; all 1s is largest
  - bias of 127 for single precision and 1023 for double precision
  - summary:  $(-1)^{\text{sign}} \times (1 + \text{significand}) \times 2^{\text{exponent} - \text{bias}}$

- Example 1 (Encoding):

For a given decimal number, construct its FP-representation

– decimal: - 0.75


→ binary: - 0.11 = - 1.1 × 2<sup>-1</sup> (normalized)

→ Floating-point exponent = - 1 + 127 = 126 = 01111110

**IEEE single precision:**

- 0.75 =    1   01111110   100000000000000000000000

# Decoding a floating-point number



- Sign indicated by first bit  $S \rightarrow (-1)^S$
- Subtract 127 from biased exponent  $E$  to obtain the actual exponent  $e = E - 127$
- Number in binary =  $(-1)^S 1.M \times 2^e$

# Example (Decoding)

0 1000 0000 1000 0000 0000 0000 0000 000

- Sign bit is zero:  
Number is positive
- Biased exponent  $E = 1000\ 0000|_2 = 128$ ;
- Actual exponent  $e = E - 127 = 1 \rightarrow 2^1$
- Significand  $\rightarrow 1.1$  (restored the hidden bit)
- The number  $= + 1.1 \times 2^1|_2 = + 11|_2 = + 3.0|_{10}$

# Example (Decoding)

[illegible]

- Sign bit is one:  
Number is negative
- Biased exponent  $E = 0111\ 1110|_2 = 126$ ;
- Actual exponent  $e = E - 127 = -1 \rightarrow 2^{-1}$
- Significand  $\rightarrow 1.11$  (restored the hidden bit)
- The number =  $-1.11 \times 2^{-1}|_2 = -0.111|_2 = -0.875|_{10}$

# Example (Encoding)

- Represent  $-2$  in FP-format
  - Convert to binary:  $10$
  - Normalize:  $1.0 \times 2^1$
  - Sign bit is 1
  - FP-exponent is  $127 + 1 = 128 = 10000000_{\text{two}}$
  - Mantissa is  $00\dots0$

1	1000 0000	0000 0000 0000 0000 0000 0000
---	-----------	-------------------------------

0/1	0000 0000 0000 0000 0000 0000 0000 0000	exponent
0	0	Zero
0	Nonzero	<p>* Any non-zero number that is smaller than the smallest normalized FP-number is a denormal number (consider magnitude only)</p>
1-254	Anything	
255	0	
255	Nonzero	NaN like 0/0



# Single-Precision Range

- Exponents 00000000 and 11111111 reserved
- Smallest value
  - Exponent: 00000001  
 $\Rightarrow$  actual exponent =  $1 - 127 = -126$
  - Mantissa: 000...00  $\Rightarrow$  normalized significand = 1.0
  - $\pm 1.0 \times 2^{-126} \approx \pm 1.2 \times 10^{-38}$
- Largest value
  - exponent: 11111110  
 $\Rightarrow$  actual exponent =  $254 - 127 = +127$
  - Mantissa: 111...11  $\Rightarrow$  normalized significand  $\approx 2.0$
  - $\pm 2.0 \times 2^{+127} \approx \pm 3.4 \times 10^{+38}$

# Double-Precision Range

- Exponents 0000...00 and 1111...11 reserved
- Smallest value
  - Exponent: 000000000001  
 $\Rightarrow$  actual exponent =  $1 - 1023 = -1022$
  - Mantissa : 000...00  $\Rightarrow$  norm. significand = 1.0
  - $\pm 1.0 \times 2^{-1022} \approx \pm 2.2 \times 10^{-308}$
- Largest value
  - Exponent: 111111111110  
 $\Rightarrow$  actual exponent =  $2046 - 1023 = +1023$
  - Mantissa: 111...11  $\Rightarrow$  norm significand  $\approx 2.0$
  - $\pm 2.0 \times 2^{+1023} \approx \pm 1.8 \times 10^{+308}$

# Floating-Point Precision

- Relative precision
  - all mantissa-bits are significant
  - Single: approx  $2^{-23}$ 
    - Equivalent to  $23 \times \log_{10} 2 \approx 23 \times 0.3 \approx 6$  decimal digits of precision
  - Double: approx  $2^{-52}$ 
    - Equivalent to  $52 \times \log_{10} 2 \approx 52 \times 0.3 \approx 16$  decimal digits of precision

# Single-Precision Normalized Range

- Exponents 00000000 and 11111111 reserved
- Smallest value


0/1 0000 0001 0000 0000 0000 0000 0000 000

- $\pm 1.0 \times 2^{-126} \approx \pm 1.2 \times 10^{-38}$

- Largest value

0/1 1111 1110 1111 1111 1111 1111 1111 111

- $= \pm(2.0 - 2^{-23}) \times 2^{+127} \approx \pm 3.4 \times 10^{+38}$



# Single-Precision FP-Denormal Numbers

- Exponents 00000000 and 11111111 reserved
- Smallest value (normalized)


0/1 0000 0001 0000 0000 0000 0000 0000 0000

■  $\pm 1.0 \times 2^{-126} \approx \pm 1.2 \times 10^{-38}$

- Denormal numbers (gradual underflow)

0/1 0000 0000 M  $\neq 0$

$= \pm 0.M \times 2^{-126}$  {min:  $\pm 2^{-23} \times 2^{-126}$ ; max:  $\pm (1-2^{-23}) \times 2^{-126}$ }



# Denormal FP-Numbers

- Exponent = 000...0  $\Rightarrow$  hidden bit is 0

$$N = (-1)^S \times (0.\text{Mantissa}) \times 2^{-126}$$

- Smaller than normalized numbers
  - allow for gradual underflow, with diminishing precision
- Denormal with mantissa = 000...0
- $\rightarrow +0, -0$

# Infinites and NaNs

- Exponent = 111...1, Mantissa = 000...0
  - $\pm$  Infinity
  - Can be used in subsequent calculations, avoiding need for **overflow check**
- Exponent = 111...1, Mantissa  $\neq$  000...0
  - Not-a-Number (NaN)
  - Indicates illegal or undefined result
    - e.g.,  $0.0 / 0.0$ ,  $\sqrt{-3}$
  - Can be used in subsequent calculations

# Floating Point Complexities

- Operations are somewhat more complicated
- In addition to overflow, we may have “underflow”
- Accuracy can be a big problem
  - IEEE 754 keeps three extra bits, guard, round, and sticky
  - several rounding modes
  - Non-zero number divide-by-zero yields “infinity” → overflow
  - Non-zero number divide-by-infinity yields → underflow
  - zero divide-by-zero yields “not a number (NaN)”
- Implementing the standard can be tricky
- Not using the standard can be even worse
- Remember the 1994 Pentium FDIV bug; write-off cost US\$ 300 M



CS 31007

Autumn 2021

# COMPUTER ORGANIZATION AND ARCHITECTURE

---

Instructors

Rajat Subhra Chakraborty (*RSC*)

Bhargab B. Bhattacharya (*BBB*)

Lecture #25: Tutorial on Floating-Point Arithmetic

23 September 2021

---

Indian Institute of Technology Kharagpur  
*Computer Science and Engineering*

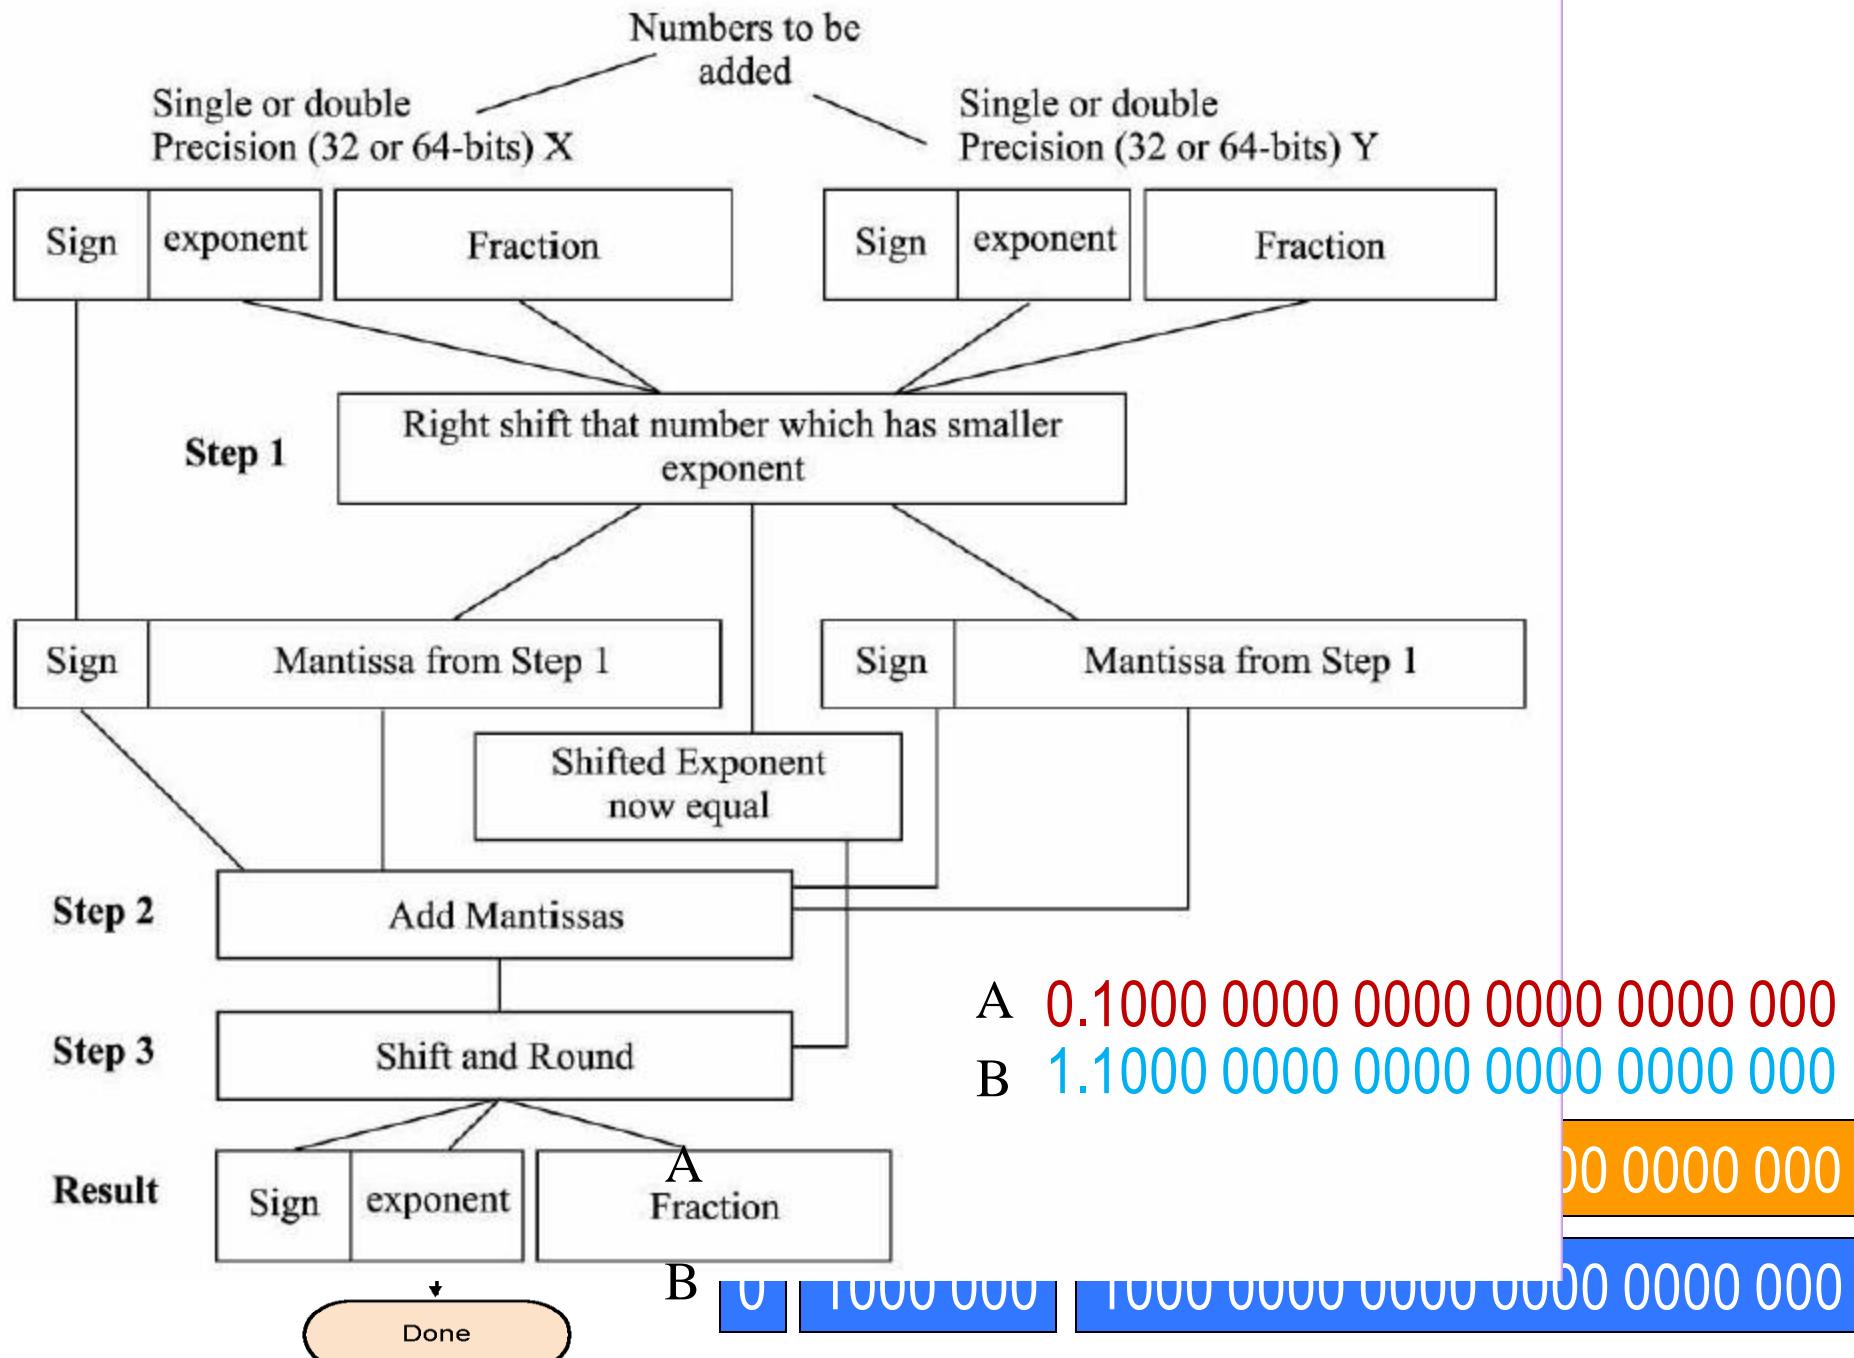
# Floating-Point Addition

overflow?

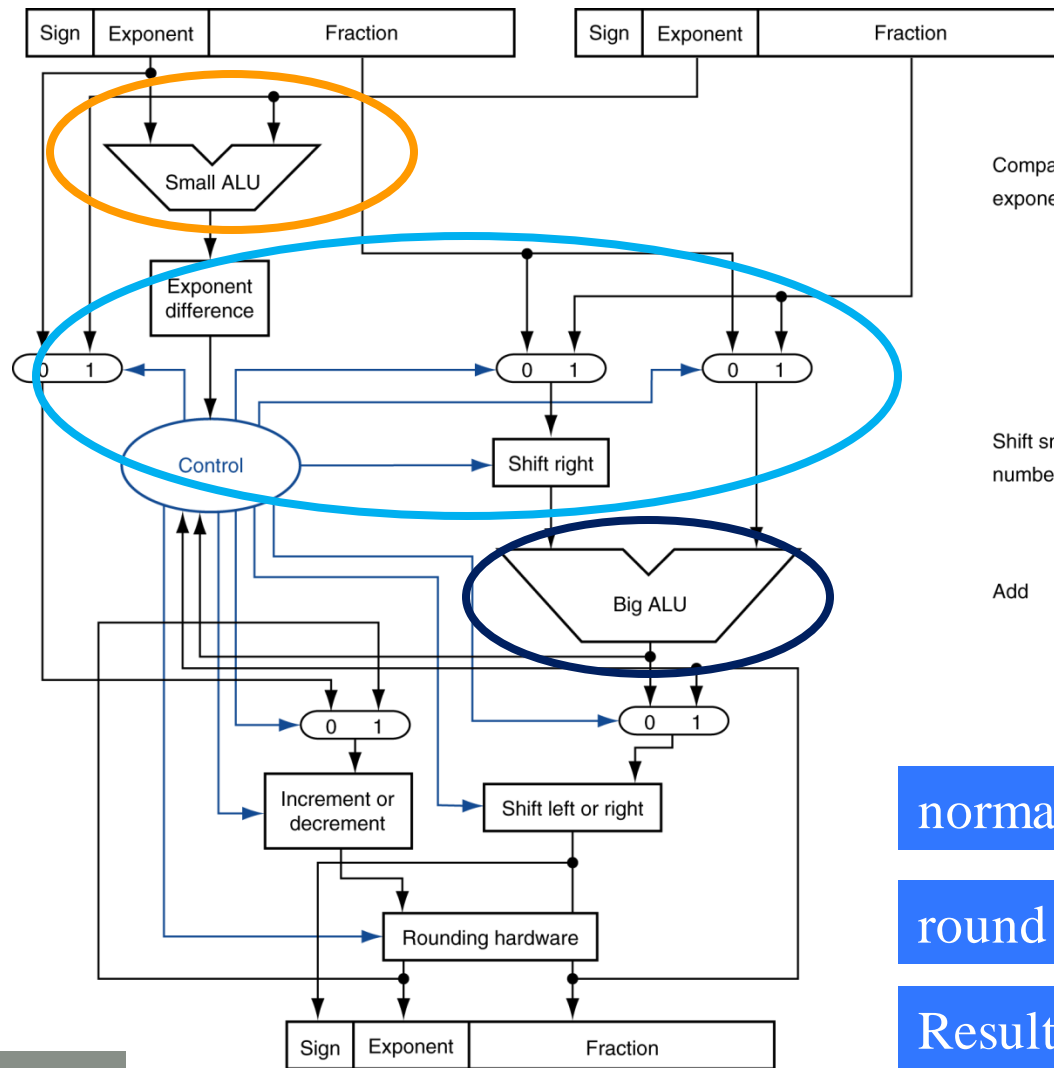
```
1.111
1.111
+ ----
11.110
```

- Consider a 3-digit mantissa binary example
  - $1.000_2 \times 2^{-1} + -1.111_2 \times 2^{-2}$
- 1. Align binary points
  - Shift **right** the number with **smaller** exponent
  - $1.000_2 \times 2^{-1} + -0.111_2 \times 2^{-1}$  ; **rightmost 1 is lost**
- 2. Add significands (integer addition)
  - $1.000_2 \times 2^{-1} + -0.111_2 \times 2^{-1} = 0.001_2 \times 2^{-1}$
- 3. Normalize result & check for over/underflow
  - $1.000_2 \times 2^{-4}$ , with no over/underflow
- 4. Round and renormalize if necessary
  - $1.000_2 \times 2^{-4}$  (no change) = 0.0625

How do  
you  
determine  
the sign of  
the result?



# FP Adder Hardware



Step 1

compare  
exponents

Step 2

equalizing  
exponents

Integer  
addition

Step 3

normalize

Step 4


round

Step 5

Result

# Three Extra Bits for Internal Use

- ❖ Three extra bits are added called **Guard, Round, Sticky**
  - ✧ To achieve accurate arithmetic such as rounding of significand, otherwise some bits would have been lost during right shifts
  - ✧ Reduces hardware without compromising precision



# Guard Bit

- ❖ When we shift bits to the right for alignment, some bits are lost
- ❖ We may need to shift the result to the left for normalization after operation
- ❖ Storing the lost bits shifted to the right will make results more accurate during normalization
- ❖ Round and Sticky bits provide further handles for accurate rounding

# Guard, Round, and Sticky Bits

- ❖ Two extra bits are needed for rounding
  - ✧ Rounding performed after **normalizing** a result significand
  - ✧ **Round bit**: appears after the guard bit
  - ✧ **Sticky bit**: appears after the round bit (**OR of all additional bits**)

$$\begin{array}{r}
 1.00000000000000000000000000000000 \\
 -0.01111111111111111111111111111111 \\
 \hline
 0.10000000000000000000000000000001 \\
 1.00000000000000000000000000000000 \\
 \hline
 0.11111111111111111111111111111110 \\
 1.11111111111111111111111111111110
 \end{array}$$

The diagram illustrates the rounding process for the subtraction of two floating-point numbers. The result of the subtraction is  $0.10000000000000000000000000000001$ . This result is then normalized by shifting it right by 7 bits, resulting in  $1.00000000000000000000000000000000$ . The normalized result is then rounded to the nearest integer, resulting in  $1.11111111111111111111111111111110$ .

The diagram also shows the intermediate steps of the rounding process. The result of the subtraction is  $0.10000000000000000000000000000001$ . This result is then normalized by shifting it right by 7 bits, resulting in  $1.00000000000000000000000000000000$ . The normalized result is then rounded to the nearest integer, resulting in  $1.11111111111111111111111111111110$ .

The diagram also shows the intermediate steps of the rounding process. The result of the subtraction is  $0.10000000000000000000000000000001$ . This result is then normalized by shifting it right by 7 bits, resulting in  $1.00000000000000000000000000000000$ . The normalized result is then rounded to the nearest integer, resulting in  $1.11111111111111111111111111111110$ .

# Floating-Point Multiplication

- Now consider a 3-digit mantissa binary example
  - $1.000_2 \times 2^{-1} \times -1.110_2 \times 2^{-2}$  (0.5 × −0.4375)
- 1. Add exponents
  - Unbiased:  $-1 + -2 = -3$
  - Biased:  $(-1 + 127) + (-2 + 127) = -3 + 254 - 127 = -3 + 127$
- 2. Multiply significands
  - $1.000_2 \times 1.110_2 = 1.110_2 \Rightarrow 1.110_2 \times 2^{-3}$
- 3. Normalize result & check for over/underflow
  - $1.110_2 \times 2^{-3}$  (no change) with no over/underflow
- 4. Round and renormalize if necessary
  - $1.110_2 \times 2^{-3}$  (no change)
- 5. Determine sign: +ve × −ve  $\Rightarrow$  −ve
  - $-1.110_2 \times 2^{-3} = -0.21875$

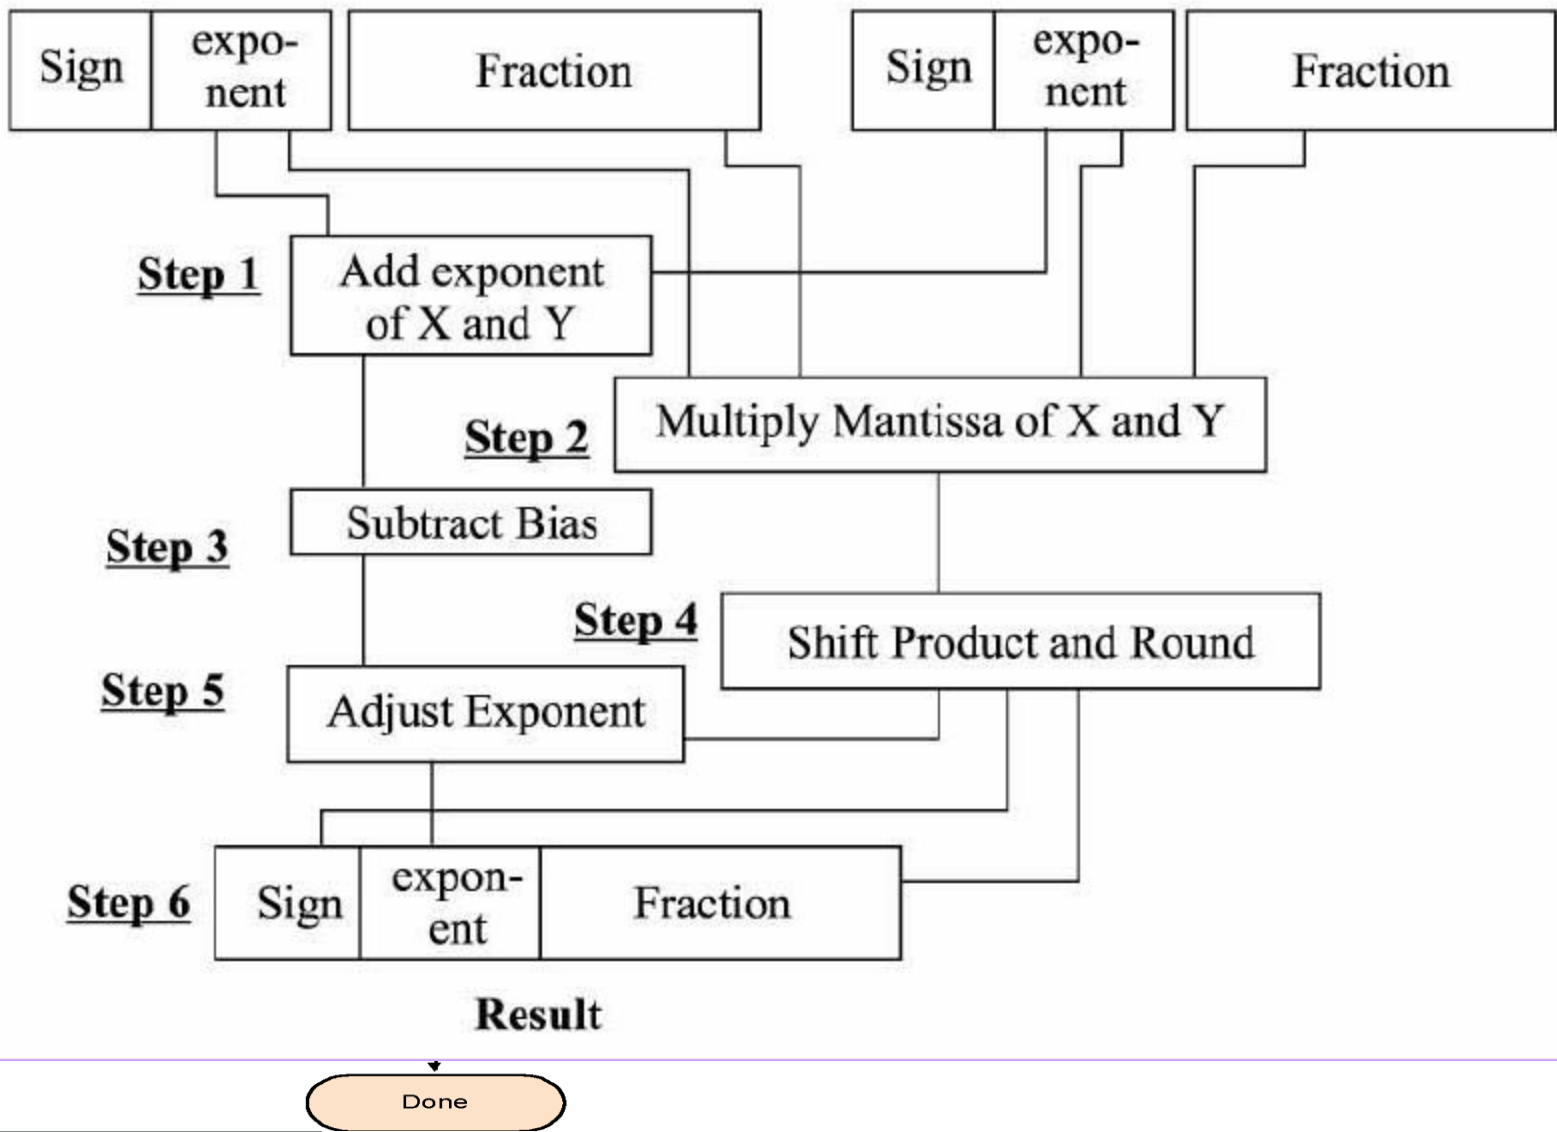


Single or double precision (32 or 64-bits) X

Numbers to be Multiplied

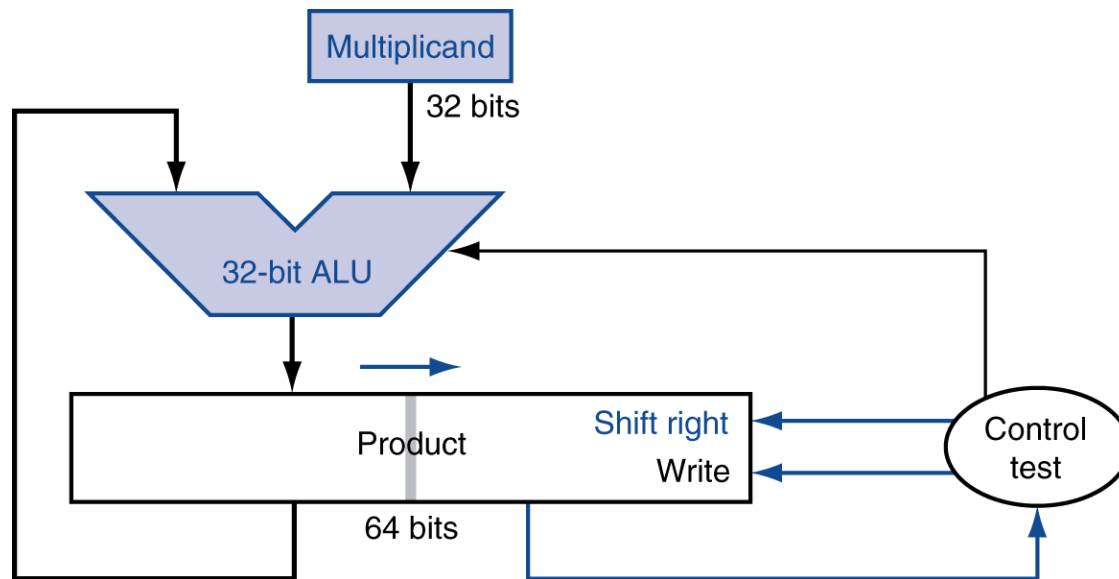
Single or double Precision (32 or 64-bits) Y

on



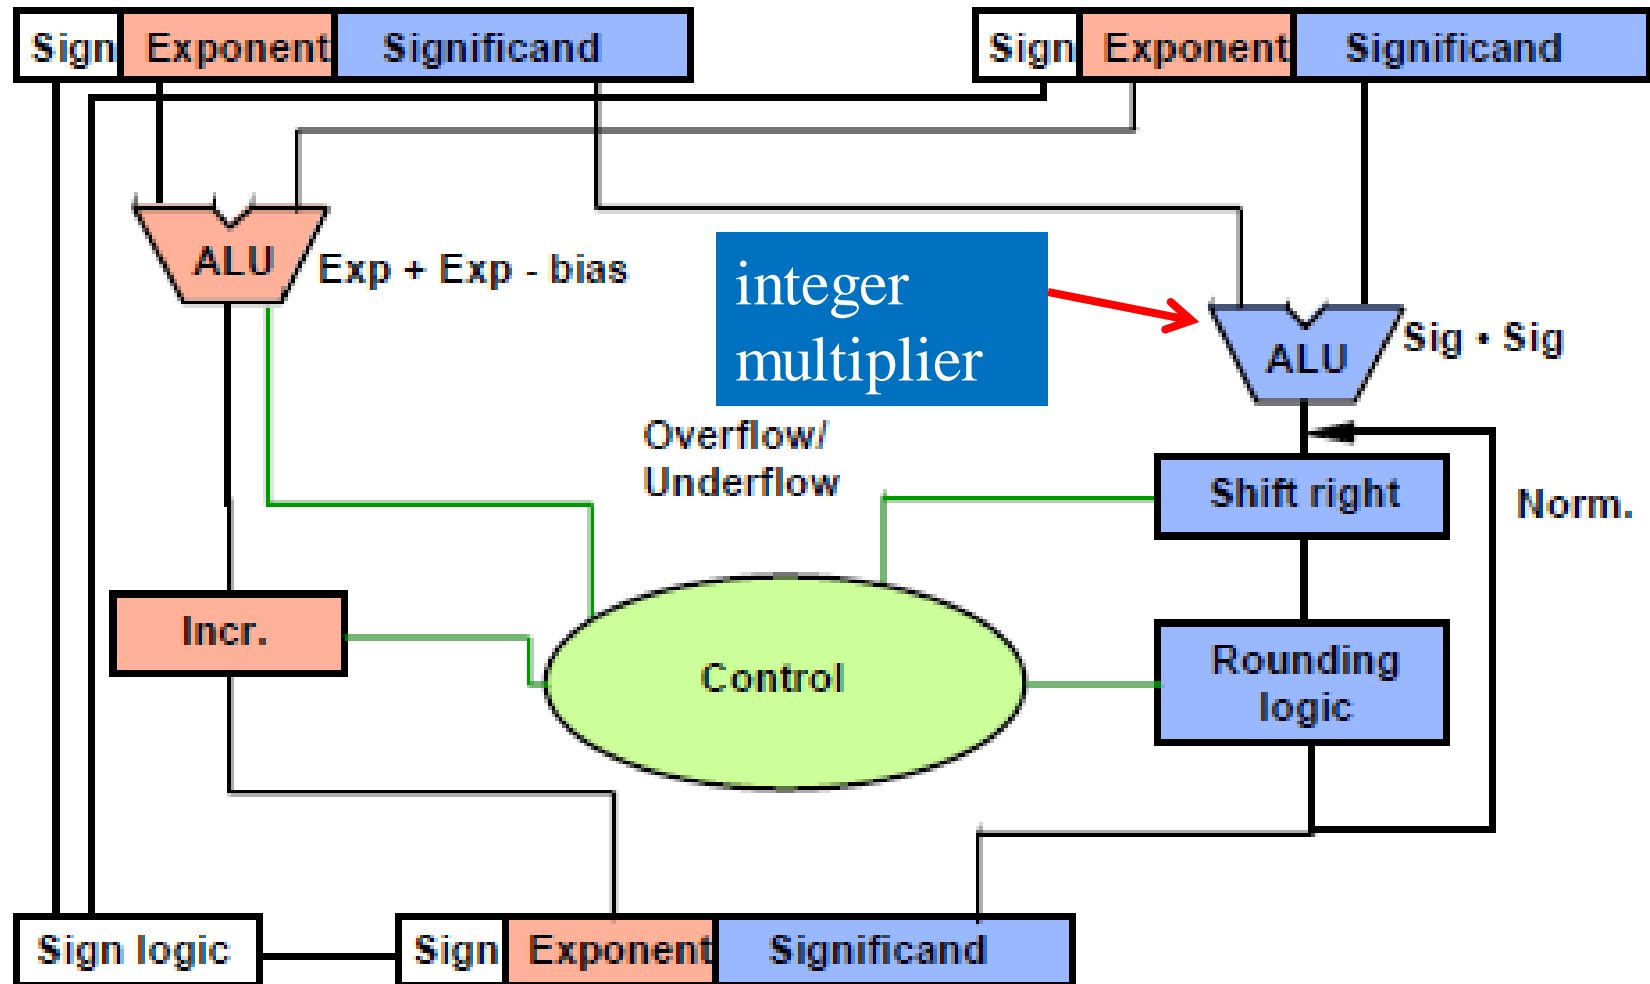
# Optimized Integer Multiplier

- Perform steps in parallel: add/shift



- One cycle per partial-product addition
  - That's ok, if frequency of multiplications is low

# FP Multiplier – Hardware Realization



# FP Arithmetic Hardware

- FP multiplier is of similar complexity to FP adder
  - But uses a multiplier for significands instead of an adder
- FP arithmetic hardware usually does
  - Addition, subtraction, multiplication, division, reciprocal, square-root
  - $FP \leftrightarrow$  integer conversion
- Operations usually takes several cycles
  - Can be pipelined

# FP Instructions in MIPS

- FP hardware is **Co-processor 1**
  - Adjunct processor that extends the ISA
- **Separate FP registers**
  - 32 single-precision: \$f0, \$f1, ... \$f31
  - Paired for double-precision: \$f0/\$f1, \$f2/\$f3, ...
    - Release 2 of MIPS ISA supports 32 × 64-bit FP reg's
- FP instructions operate only on FP registers
  - Programs generally don't do integer ops on FP data, or vice versa
  - More registers with minimal code-size impact
- FP load and store instructions
  - lwc1, swc1,
    - e.g., lwc1 \$f8, 32(\$sp)

# FP Instructions in MIPS

- Single-precision arithmetic
  - `add.s`, `sub.s`, `mul.s`, `div.s`
    - e.g., `add.s $f0, $f1, $f6`
- Double-precision arithmetic
  - `add.d`, `sub.d`, `mul.d`, `div.d`
    - e.g., `mul.d $f4, $f4, $f6`
- Single- and double-precision comparison
- Branch on FP condition code true or false

# FP Example: °F to °C

- C code:

```
float f2c (float fahr) {  
    return ((5.0/9.0)*(fahr - 32.0));  
}
```

- fahr in \$f12, result in \$f0, literals in global memory space

- Compiled MIPS code:

```
f2c: lwc1    $f16, const5($gp)  
     lwc2    $f18, const9($gp)  
     div.s   $f16, $f16, $f18  
     lwc1    $f18, const32($gp)  
     sub.s   $f18, $f12, $f18  
     mul.s   $f0, $f16, $f18  
     jr      $ra
```

# Summary

- Computer arithmetic is constrained by limited precision
- Bit patterns have no inherent meaning but standards do exist
  - two's complement
  - IEEE 754 floating point
- Computer instructions determine “meaning” of the bit patterns
- Performance and accuracy are important so there are many complexities and implementation challenges in real machines



*Problem:* Compute  $z = \lfloor s/9 \rfloor$  using only integer addition/subtraction, and shift;  $s$ : integer  
Division operation not allowed

$$z = \frac{s}{9}$$

Hardware requirement:  
Only shifter and integer adder;  
Iterate a few times to obtain a  
close approximation!

The diagram illustrates the iterative process of calculating  $z = \lfloor s/9 \rfloor$  using only integer addition, subtraction, and shifts. A red box highlights the iterative step. A blue line traces the sequence of operations from the initial value  $z$  to the final result.

$$\bar{z} = \frac{s - \frac{s - \frac{s - \dots}{8}}{8}}{8} \quad \text{to } \infty$$

$$\begin{aligned} z &= (s - z)/8 \\ \Rightarrow 8z &= s - z \\ \Rightarrow z &= s/9 \end{aligned}$$